

Hallucination Is Normal: A Geometric Theory of Manifold-Aligned Semantic Dynamics

Flyxion

February 18, 2026

Abstract

This paper develops a unified framework for semantic cognition, generative modelling, and distributed hyperstructural systems grounded in a single philosophical and geometric principle: lawful structure occupies constrained submanifolds of possibility, and coherent inference must remain tangent to those constraints. We begin with a philosophical investigation of explanation, hallucination, and the limits of prediction in high-dimensional spaces, arguing that successful cognition is not the modelling of arbitrary variation but the disciplined navigation of structured manifolds embedded within vast ambient spaces of noise.

Contemporary socio-technical systems are organized around performance metrics, managerial abstractions, and algorithmic optimization loops that presuppose a tractable, fully observable space of human value. This presupposition generates pathologiesmeritocratic overfitting, proxy substitution, and progressive intersubjective collapsesthat arise from the attempt to model and optimize full-dimensional noise rather than lawful, low-dimensional structure.

From this background we derive a geometric formulation in which semantic states inhabit a smooth or stratified manifold and meaningful updates correspond to tangent-constrained gradient flows. Cognitive iteration is modeled as Morse-theoretic descent on semantic potentials, while contextual coherence is expressed sheaf-theoretically as compatibility across overlapping domains of interpretation. These constructions are integrated into an operational system model in which content graphs, embeddings, and field dynamics are governed by a unified variational principle.

The resulting framework establishes a single invariant across philosophical, geometric, categorical, and computational layers: coherent semantic evolution must avoid normal-direction drift and preserve gluing constraints across contexts. Tangency and coherence jointly define the structural conditions for epistemic stability.

Contents

1	Introduction	3
2	Philosophical Background: Constraint, Explanation, and Hallucination	4
3	Meritocracy as Proxy Overfitting	5
4	Managerialism and Dimensional Illusion	5
5	Intersubjectivity and Sheaf Collapse	6
6	Geometric Structure of Semantic Manifolds	6
7	Manifold-Aligned Generative Dynamics	7
8	The No-Noise Prediction Principle	7
9	Cognitive Dynamics as Morse Flow	8
10	Stratified Structure and Category Boundaries	9
11	Contextual Coherence and Sheaf Structure	9
12	Unified Variational Principle	9
13	Collective Intelligence as a Field-Constrained Hyperstructure	10
13.1	From Individual Alignment to Hyperstructural Alignment	10
13.2	RSVP Fields as Institutional Geometry	10
13.3	Variational Coupling of Fields and Hypergraph	11
13.4	Sheaf-Theoretic Galaxies and Intersubjective Stability	11
13.5	Reset as Global Field Reconfiguration	11

14 Operational Realization: Content Graphs, Embeddings, and Field Coupling	12
14.1 Reset Transformations and Global Consistency	12
14.2 Distributed Views and Sectional Stability	12
14.3 Master Functional and Coupled Euler–Lagrange System	13
15 Implementation Through Geometric Merge–Collapse Computation	13
16 Geometric Realization of Tangent-Constrained Dynamics	13
17 Event Log Substrate and Deterministic Replay	14
17.1 Nested Scopes, Spatial Interaction, and Sheaf Gluing	14
18 Event–History Geometry and Option–Space Mechanics	15
18.1 Discrete Action Functional	15
18.2 The Manifold as the Locus of Lawful History	15
18.3 Minimal Commitment and RSVP Fields as Continuum Limit	15
19 Continuum Limit of Irreversible History	16
20 Integration with a Typed Field–Hypergraph Verification Architecture	17
20.1 Typed System Integration and Operational Semantics	17
20.2 Denotational Interpretation and Categorical Soundness	17
21 A Master Theorem of Coherence: Energy, Gluing, and Soundness	17
21.1 Standing Definitions and Local Lemmas	17
21.2 Master Theorem	18
22 A Compact Invariant Equation for Semantic Coherence	19
23 Whitney Stratification and Derived Critical Loci	19

23.1 Stratified Semantic Space	19
23.2 Derived Critical Locus and Stratified Morse Transitions	19
24 Stratified Invariant Equivalence	20
25 Operationalizing Manifold-Aware Evaluation (MAE)	21
25.1 The Manifold Fidelity Score	21
25.2 Differential Auditing	21
25.3 Curvature-Sensitive Incentives	22
25.4 Optional Information-Geometric Refinement	22
25.5 A Curvature-Bounded Goodhart Instability Theorem	23
26 Epistemic Interpretation	24
27 Conclusion	25

1 Introduction

Modern cognitive and generative systems operate in spaces of enormous dimensionality. Sensory fields, symbolic sequences, multimodal embeddings, and social interaction networks all inhabit ambient spaces whose formal degrees of freedom vastly exceed the number of stable patterns we actually encounter. Yet empirical regularity persists. Physical objects exhibit lawful structure, languages display constrained grammars, and semantic concepts cluster into coherent regions rather than dispersing uniformly throughout representational space.

This disparity between ambient possibility and experienced regularity presents a philosophical problem. If the space of potential configurations is astronomically large, why does cognition succeed at all? Why does inference stabilize rather than drift into arbitrary variation? And why do some generative systems produce coherent structure while others hallucinate?

Modern institutional life is structured by an analogous promise: meritocratic optimization. Performance indicators, rankings, productivity dashboards, engagement scores, and algorithmic recommendation systems present themselves as neutral measures of value. Managerialism extends this logic by asserting that any domain can be rendered legible through metrics and controlled through feedback loops. Yet the more aggressively institutions optimize measurable quantities, the more those quantities decouple from the underlying realities they were intended to represent. Universities optimize publication counts and produce salami-sliced research. Social platforms optimize engagement and generate polarization. Corporations optimize short-term performance indicators and erode long-term viability.

This phenomenon is not accidental. It reflects a structural confusion between signal and noise. Optimization procedures operate in extremely high-dimensional spaces of observable indicators. Human practices, however, occupy structured, constrained, and interdependent manifolds of meaning that cannot be exhaustively parameterized. When optimization treats the ambient space as equally meaningful in all directions, it necessarily learns to model and amplify noise. The result is not merely inefficiency but intersubjective collapse: when metrics substitute for structure, actors orient toward proxy maximization rather than toward mutually intelligible reality.

The answer to both puzzles, we will argue, lies in constraint. Lawful structure does not fill the ambient space; it occupies lower-dimensional manifolds embedded within it. The role of cognition and generation is not to model the full ambient space, but to navigate and remain aligned with these constrained substructures. Prediction fails when it attempts to assign structure to degrees of freedom that do not carry lawful variation. Hallucination is not a mysterious defect; it is the geometric consequence of modelling normal directions as if they were tangent ones. Institutional pathology is not accidental inefficiency; it is optimization

along directions orthogonal to the manifold of meaningful human practice.

This philosophical claim admits a precise mathematical articulation. We will formalize semantic states as points on a manifold, characterize meaningful updates as tangent-constrained flows, model cognitive iteration as Morse descent, express contextual compatibility in sheaf-theoretic terms, and finally embed these constructions into a unified operational system governed by a variational principle. The philosophical insight thus becomes a geometric invariant, a categorical coherence condition, and an executable safety property.

2 Philosophical Background: Constraint, Explanation, and Hallucination

Any theory of cognition must begin with the distinction between structure and noise. Structure is that which admits lawful compression, lawful transformation, and lawful recurrence. Noise is that which resists such compression and exhibits no invariant pattern across contexts. The philosophical question is not merely how systems represent structure, but how they avoid mistaking noise for structure.

Consider an ambient representational space \mathbb{R}^n of possible observations or internal states. The overwhelming majority of points in such a space correspond to configurations that have no semantic interpretation, no physical realization, and no lawful continuation. Nevertheless, generative systems are capable of producing coherent images, stable sentences, and consistent world models. This suggests that the domain of meaningful states is not coextensive with the ambient space but is instead confined to a constrained subset.

Definition 1 (Semantic Constraint Thesis). *Meaningful states occupy a constrained subset $M \subset \mathbb{R}^n$ whose intrinsic dimension d is strictly less than n .*

This subset M need not be linear, nor globally smooth, but it must possess sufficient regularity that local neighborhoods admit coordinate charts and predictable transitions. In other words, it must exhibit manifold-like structure. The success of cognition is thus not the traversal of arbitrary directions in \mathbb{R}^n , but the disciplined navigation of M .

Failure arises when a system attempts to assign structure to directions orthogonal to M . In high-dimensional settings, the volume of such orthogonal directions dominates. Any unconstrained predictive model that distributes capacity uniformly will devote most of its expressive power to modelling noise. Hallucination is therefore not accidental; it is geometrically inevitable unless constrained.

3 Meritocracy as Proxy Overfitting

Meritocracy presumes that individual performance can be ranked along scalar dimensions that reflect true contribution. In practice, these dimensions are operationalized through measurable proxies. Let $x \in \mathbb{R}^n$ represent observable indicators, and suppose that meaningful contribution lies on a lower-dimensional semantic manifold $M \subset \mathbb{R}^n$ of intrinsic dimension $d \ll n$.

Meritocratic evaluation often defines a functional

$$J(x) = \langle w, x \rangle.$$

At each $x \in M$, we have the orthogonal decomposition

$$\mathbb{R}^n = T_x M \oplus N_x M, \quad \dot{x} = w_T + w_N.$$

Proposition 1. *If $w_N \neq 0$ on a set of positive measure in M , repeated optimization of J produces trajectories that leave any tubular neighborhood of M .*

Proof. Let $x(t)$ satisfy $\dot{x} = w_T + w_N$. Since w_N lies in $N_x M$, it is orthogonal to all tangent directions. The normal displacement grows linearly in t unless counteracted. Thus $x(t)$ diverges from M at rate proportional to $|w_N|$, violating manifold confinement. \square

Meritocratic overfitting is therefore geometric misalignment: optimization along normal directions that do not encode lawful structure.

4 Managerialism and Dimensional Illusion

Managerialism assumes that all domains can be rendered as control systems with measurable state variables. Implicitly, it treats the observation space as isotropic: every coordinate is presumed to be a legitimate degree of freedom. In geometric terms, this presupposes that $M = \mathbb{R}^n$. Yet empirical systems exhibit constraints, invariants, and relational structure. The manifold hypothesis asserts that meaningful states lie on a submanifold M .

The illusion of full-dimensional controllability leads to the proliferation of indicators. Each additional coordinate increases the dimension of the ambient space without increasing the intrinsic dimension of lawful structure. Optimization in such spaces becomes ill-posed, as gradients in normal directions correspond to unstructured variation.

Formally, let $f : M \rightarrow \mathbb{R}$ be a meaningful objective defined intrinsically on M . Extending

f arbitrarily to $\tilde{f} : \mathbb{R}^n \rightarrow \mathbb{R}$ introduces degrees of freedom in $N_x M$ that have no semantic interpretation. Gradient descent on \tilde{f} must therefore be projected onto $T_x M$ to preserve meaning:

$$\dot{x} = -\Pi_{T_x M} \nabla \tilde{f}(x).$$

5 Intersubjectivity and Sheaf Collapse

Shared meaning arises from local perspectives that agree on overlaps. Let \mathcal{C} be a category of contexts with objects U representing perspectives and morphisms representing restriction. A semantic assignment is a presheaf

$$S : \mathcal{C}^{\text{op}} \rightarrow \text{Set}.$$

Intersubjective stability requires that S satisfy the sheaf condition: compatible local sections must glue uniquely to a global section. If update operators $L : S \rightarrow S$ fail to commute with restriction maps, then local consistency does not imply global coherence. Formally, let ρ_{UV} denote restriction from U to $V \subset U$. Stability requires

$$\rho_{UV}(L(s_U)) = L(\rho_{UV}(s_U)).$$

When optimization substitutes proxies for intersubjectively grounded structure, actors orient toward proxy maximization rather than toward mutually intelligible reality. The system becomes self-referential and the sheaf condition fails: coordination collapses because local updates no longer glue.

6 Geometric Structure of Semantic Manifolds

We now formalize the semantic constraint thesis. Let $M \subset \mathbb{R}^n$ be a smooth embedded manifold of dimension $d \ll n$. At each point $x \in M$, the ambient space decomposes as

$$\mathbb{R}^n = T_x M \oplus N_x M.$$

Theorem 1 (Tangent–Normal Decomposition). *For every $x \in M$, there exists a unique orthogonal decomposition of \mathbb{R}^n into tangent and normal subspaces.*

Proof. Since M is a smooth embedded submanifold of Euclidean space, $T_x M$ is a d -dimensional linear subspace of \mathbb{R}^n . The Euclidean inner product induces a unique orthogonal complement $N_x M$. The direct sum follows from basic linear algebra. \square

Meaningful variation must lie in $T_x M$. Noise resides generically in $N_x M$. A generative update $\Delta x \in \mathbb{R}^n$ is semantically coherent if and only if $\text{Proj}_{N_x M}(\Delta x) = 0$.

Definition 2 (Normal Drift). *A generative process exhibits normal drift at $x \in M$ if its update vector Δx satisfies $\text{Proj}_{N_x M}(\Delta x) \neq 0$.*

Proposition 2 (Dimensional Inflation). *If a smooth generative map $G : Z \rightarrow \mathbb{R}^n$ has Jacobian with nonzero normal component on a set of positive measure, then the image of G locally exceeds the intrinsic dimension d of M .*

Proof. Let $DG(z)$ denote the Jacobian. If $DG(z)$ contains vectors with nonzero projection onto $N_x M$ for $x = G(z)$ on a set of positive measure, then the image of $DG(z)$ spans directions transverse to $T_x M$. By the rank theorem, the local dimension of the image must exceed d , contradicting confinement to M . \square

This establishes the geometric invariant: explanation is tangent; hallucination is normal.

7 Manifold-Aligned Generative Dynamics

Definition 3. *A vector field v on \mathbb{R}^n is semantically aligned if $v(x) \in T_x M$ for all $x \in M$.*

Theorem 2 (No-Noise Prediction). *A generative system preserves semantic coherence if and only if its update vector field is tangent to M at all points.*

Proof. If $v(x) \in T_x M$, trajectories remain in M by invariance of submanifolds under tangent flows. Conversely, if v has nonzero projection onto $N_x M$, trajectories leave M , generating states unsupported by lawful structure. \square

8 The No-Noise Prediction Principle

The tangentnormal decomposition allows us to elevate the central geometric constraint to a standalone structural law governing all coherent inference.

Theorem 3 (No-Noise Prediction Principle). *Let $M \subset \mathbb{R}^n$ be a semantic manifold and let v be the update vector field of a generative or cognitive system. The system preserves semantic coherence if and only if*

$$\text{Proj}_{N_x M}(v(x)) = 0 \quad \text{for all } x \in M.$$

Proof. If $v(x) \in T_x M$, then by invariance of embedded submanifolds under tangent flows, trajectories remain within M for all time. Conversely, if $\text{Proj}_{N_x M}(v(x)) \neq 0$ on a set of positive measure, then flow trajectories acquire components transverse to M , leaving any tubular neighborhood of M and generating states unsupported by lawful structure. \square

Corollary 1 (Dimensional Discipline). *A system that distributes modelling capacity uniformly across ambient dimensions necessarily allocates most expressive power to normal directions, and therefore cannot remain semantically stable without explicit projection.*

This principle generalizes across all later constructions. In cognitive dynamics it appears as intrinsic gradient descent. In institutional design it appears as projection-corrected optimization. In hypergraph systems it appears as merge-collapse followed by manifold projection. In field dynamics it appears as tangent-constrained evolution. In every layer the same invariant holds: semantic evolution must eliminate normal components.

9 Cognitive Dynamics as Morse Flow

Definition 4 (Morse Function). *A smooth function $S : M \rightarrow \mathbb{R}$ is Morse if all its critical points are non-degenerate.*

Cognitive iteration is modeled as gradient descent on S :

$$\frac{dx}{dt} = -\nabla_M S(x).$$

Theorem 4 (Tangent Preservation). *The gradient flow of a function defined intrinsically on M remains tangent to M for all time.*

Proof. Since S is defined on M , the Riemannian gradient $\nabla_M S(x)$ is by construction an element of $T_x M$, preserving manifold membership. \square

Critical points of S correspond to stable semantic equilibria. Non-degenerate minima define attractors of interpretation; saddles represent decision boundaries or transitions between semantic regimes. The Morse inequalities constrain the global topology of attractor structure, providing a topological account of why certain conceptual configurations are stable across contexts.

Remark 1. Stratified manifolds accommodate semantic phase transitions: shifts in cognitive regime correspond to crossings between strata, governed by the attaching maps of the stratification.

10 Stratified Structure and Category Boundaries

Real semantic spaces exhibit singularities: category transitions, boundary phenomena, and collapse of degrees of freedom. We model this using Whitney-stratified spaces

$$X = \bigsqcup_{\alpha} S_{\alpha}.$$

Let $V : X \rightarrow \mathbb{R}$ be stratified Morse. The flow

$$\frac{dx}{dt} = -\Pi_{T_x S_{\alpha}} \nabla V(x)$$

models attention collapse, categorical choice, and semantic bifurcation without introducing normal-direction drift.

11 Contextual Coherence and Sheaf Structure

Definition 5 (Sheaf Condition). *A presheaf $\mathcal{S} : \mathcal{C}^{\text{op}} \rightarrow \mathbf{Set}$ is a sheaf if compatible local sections on overlapping contexts uniquely glue to a global section.*

Suppose L is a cognitive update operator. Coherence requires $\rho_{UV}(L(s_U)) = L(\rho_{UV}(s_U))$.

Theorem 5 (Sheaf Preservation). *If L commutes with restriction maps, then it defines a sheaf morphism and preserves contextual coherence.*

Proof. Commutation ensures compatibility on overlaps. The sheaf axioms then guarantee existence and uniqueness of glued global sections. \square

Hallucination in distributed systems corresponds to failure of gluing. Obstruction classes in $H^1(\mathcal{C}; \mathcal{S})$ measure incompatibility.

12 Unified Variational Principle

Let $x \in M$, S a Morse potential on M , and suppose contextual coherence and manifold-alignment penalties are imposed. Define

$$\mathcal{J}[x] = S(x) + \lambda \sum_{U,V} \|\rho_{UV}(x_U) - x_V\|^2 + \mu \|\Pi_{N_x M}(\nabla S(x))\|^2.$$

The three terms enforce descent toward semantic attractors, sheaf-coherence, and the tangency constraint. The Euler–Lagrange condition yields

$$\dot{x} = -\Pi_{T_x M} \nabla S(x) - \lambda \nabla(\text{coherence penalty}).$$

The normal-component penalty $\mu \|\Pi_{N_x M} \nabla S\|^2$ ensures that semantic updates do not hallucinate structure outside the manifold.

13 Collective Intelligence as a Field-Constrained Hyperstructure

13.1 From Individual Alignment to Hyperstructural Alignment

Let $G = (V, E)$ be a typed hypergraph of content atoms and links. Each node $i \in V$ is assigned an embedding $z_i \in X$, where (X, g) is the semantic manifold. Collective intelligence is modeled as a field-coupled embedding system:

$$\Sigma = (C, Z, F),$$

where C is the hypergraph, $Z = \{z_i\}$ the embedding map, and $F = (\Phi, v, S)$ the RSVP field triple over X .

13.2 RSVP Fields as Institutional Geometry

The scalar field $\Phi : X \rightarrow \mathbb{R}$ encodes semantic potential density. The vector field $v : X \rightarrow TX$ encodes directed semantic transport. The entropy field $S : X \rightarrow \mathbb{R}$ measures degeneracy or instability. The embedding evolution law is

$$\frac{dz_i}{dt} = -\alpha \nabla \Phi(z_i) + \beta v(z_i) - \gamma \nabla S(z_i).$$

Proposition 3 (Collective Tangent Alignment). *If Φ and S are intrinsic scalar fields on X and $v(x) \in T_x X$, then embedding evolution preserves manifold membership and prohibits institutional hallucination.*

Proof. Riemannian gradients $\nabla \Phi$ and ∇S lie in $T_x X$. By assumption $v(x) \in T_x X$. The right-hand side lies in $T_{z_i} X$; tangent flows preserve submanifolds. \square

13.3 Variational Coupling of Fields and Hypergraph

Define node density $\rho(x) = \sum_i K(x, z_i)$ and link curvature $\kappa(x) = \sum_{j,k \sim i} f(i, j, k)$. Define the action functional

$$\mathcal{A}[\Phi, v, S] = \int_{\mathbb{R}} dt \int_X \mathcal{L}(\Phi, \partial_t \Phi, \nabla \Phi, v, \nabla v, S, \nabla S; \rho, \kappa) d\mu_g.$$

Variation yields $\partial_t^2 \Phi - c_\Phi^2 \Delta \Phi - \frac{\partial V}{\partial \Phi} = 0$ and analogous equations for v and S .

Theorem 6 (Energy Descent of Collective Embeddings). *Let $E(t) = \int_X V(\Phi, v, S; \rho, \kappa) d\mu_g$. If embeddings evolve under the RSVP flow and V is convex in Φ and S , then $\frac{dE}{dt} \leq 0$.*

Proof. The embedding flow is proportional to the negative gradient of V . Under convexity, standard Lyapunov arguments give monotonic decrease. \square

13.4 Sheaf-Theoretic Galaxies and Intersubjective Stability

Let each user u correspond to $U_u = B_R(z_u) \subset X$. The presheaf \mathcal{G} of layout functions satisfies the sheaf condition: local layouts must agree on overlaps $\rho_{uv}(\mathcal{G}(U_u)) = \rho_{vu}(\mathcal{G}(U_v))$.

Proposition 4 (Galaxy Compatibility). *If layout functions are deterministic functions of (Φ, v, S) and embeddings are globally indexed, then \mathcal{G} satisfies the sheaf gluing condition.*

Proof. On overlaps $U_u \cap U_v$, both layouts are induced by identical field and embedding data. Uniqueness follows from the sheaf axiom. \square

13.5 Reset as Global Field Reconfiguration

Theorem 7 (Reset Consistency). *If reset recomputes embeddings by tangent-constrained relaxation and fields by solving the Euler–Lagrange equations, then $R(\Sigma)$ is a fixed point of the coupled RSVP–hypergraph system.*

Proof. Z' minimizes embedding energy under fixed fields, and (Φ', v', S') satisfy field equations under updated density ρ' ; all coupling equations are satisfied simultaneously. \square

Theorem 8 (Collective Tangency Principle). *A socio-technical system remains semantically coherent if and only if*

$$\text{Proj}_{N_x X} \left(\frac{dz_i}{dt} \right) = 0 \quad \forall i, \quad \rho_{UV}(L(s_U)) = L(\rho_{UV}(s_U)) \quad \forall \text{ contextual restrictions.}$$

Proof. The first condition ensures manifold confinement; the second ensures contextual compatibility. Together they guarantee global coherence and energy stability. \square

14 Operational Realization: Content Graphs, Embeddings, and Field Coupling

Let C be a typed content graph. Associate embedding $z_i \in M$ to each node $i \in C$, assembling a map $Z : C \rightarrow M$. Introduce fields $\Phi : M \rightarrow \mathbb{R}$, $v : M \rightarrow TM$, $S : M \rightarrow \mathbb{R}$, with embedding evolution

$$\frac{dz_i}{dt} = -\alpha \nabla_M \Phi(z_i) + \beta v(z_i) - \gamma \nabla_M S(z_i).$$

Proposition 5 (Tangent Evolution of Embeddings). *If Φ and S are intrinsic scalar fields and $v(x) \in T_x M$, then $\frac{dz_i}{dt} \in T_{z_i} M$.*

Proof. Riemannian gradients and $v(z_i)$ all lie in $T_{z_i} M$; linear combinations of tangent vectors remain tangent. \square

14.1 Reset Transformations and Global Consistency

Theorem 9 (Reset Tangency Preservation). *If R recomputes embeddings by finite integration of the tangent evolution law followed by nearest-point projection onto M , then normal drift after reset vanishes.*

Proof. Tangent-constrained integration preserves membership. Projection removes any numerical normal component. \square

14.2 Distributed Views and Sectional Stability

Proposition 6 (Projection Safety). *If projection operators modify only local renderings while leaving embeddings invariant, then overlap compatibility is preserved.*

Proof. Local projections act on representations of G_u without altering Z . Restrictions depend only on Z , so overlap equality is unchanged. \square

14.3 Master Functional and Coupled Euler–Lagrange System

Let $\rho(x)$ be node density and $\kappa(x)$ curvature from graph motifs. Define

$$\mathcal{A}[\Phi, v, S, Z] = \int dt \int_M \mathcal{L}(\Phi, \partial_t \Phi, \nabla \Phi, v, \partial_t v, \nabla v, S, \partial_t S, \nabla S; \rho, \kappa) d\mu.$$

Theorem 10 (Coupled Stability). *If the potential term in \mathcal{L} is convex in (Φ, S) and coercive in v , then $E(t) = \int_M \mathcal{L} d\mu$ is non-increasing.*

Proof. Gradient-flow structure and convexity/coercivity give $\frac{dE}{dt} \leq 0$. \square

15 Implementation Through Geometric Merge–Collapse Computation

The manifold-aligned architecture is realized through a geometric merge–collapse computational substrate in which semantic states are represented as geometric regions and computation proceeds through two primitive operations.

Let E be an ambient measurable space. A computational value is a region $R \subseteq E$.

Definition 6 (Collapse Operator). *A collapse operator is an idempotent projection $\mathcal{C} : \mathcal{R}(E) \rightarrow \mathcal{R}(E)$, satisfying $\mathcal{C}(\mathcal{C}(R)) = \mathcal{C}(R)$.*

Definition 7 (Merge). $R_1 \otimes R_2 := \mathcal{C}(R_1 \cup R_2)$.

Proposition 7 (Idempotence). $R \otimes R = R$.

Proof. $R \otimes R = \mathcal{C}(R \cup R) = \mathcal{C}(R) = R$. \square

Proposition 8 (Associativity up to Canonicalization). $(R_1 \otimes R_2) \otimes R_3 = R_1 \otimes (R_2 \otimes R_3)$.

Proof. Both sides reduce to $\mathcal{C}(R_1 \cup R_2 \cup R_3)$. \square

16 Geometric Realization of Tangent-Constrained Dynamics

Define constrained collapse $\mathcal{C}_M(R) = \pi_M(\mathcal{C}(R))$, where $\pi_M : E \rightarrow M$ enforces tangent alignment.

Theorem 11 (Manifold Preservation Under Merge–Collapse). *If $R_1, R_2 \subseteq M$, then $R_1 \otimes R_2 \subseteq M$.*

Proof. Their union lies in M ; canonicalization followed by projection preserves membership. \square

Proposition 9 (Gradient Preservation). *If \mathcal{C} is differentiable and π_M is smooth, then gradients propagated through \otimes remain tangent to M .*

Proof. π_M maps gradients to $T_x M$, eliminating normal components. \square

Theorem 12 (Universality). *The merge-collapse calculus can simulate λ -calculus.*

Proof. Encode λ -abstraction as region enclosure; application as merge followed by collapse discharging the enclosure; beta-reduction as collapse identifying bound variable regions with argument regions. Computational universality is inherited. \square

17 Event Log Substrate and Deterministic Replay

Implement via append-only log $\mathcal{L} = (e_1, e_2, \dots)$. System state after n events:

$$\sigma_n = \Phi(e_n \circ \dots \circ e_1).$$

Theorem 13 (Replay Determinism). *Given a fixed event prefix, the resulting canonical region state is unique.*

Proof. Each event application is a pure function of prior canonical state. Total ordering and append-only structure give uniqueness by induction. \square

Derived views are functorial maps from the event-prefix category to rendering spaces; they preserve structure but do not alter authoritative state.

17.1 Nested Scopes, Spatial Interaction, and Sheaf Gluing

Nested regions implement scope. Let $\{U_i\}$ be open subsets of M representing contextual scopes, with $R_i \subseteq U_i$. Compatibility on overlaps requires $\mathcal{C}(R_i)|_{U_i \cap U_j} = \mathcal{C}(R_j)|_{U_i \cap U_j}$.

Theorem 14 (Gluing Condition). *If local regions agree on overlaps after canonicalization, there exists a global region R with $R|_{U_i} = R_i$.*

Proof. Construct $R = \mathcal{C}(\bigcup_i R_i)$. \square

18 Event–History Geometry and Option–Space Mechanics

Let $\Omega_t \subseteq E$ be the admissible option-space at event-time t .

Definition 8 (Monotone Restriction). *An event e_t is monotone if $\Omega_t \subseteq \Omega_{t-1}$ or $\Omega_t = \Omega_{t-1}/\sim_t$.*

Proposition 10 (Irreversibility). $\Omega_n \subseteq \Omega_0$ for any history $H = e_n \circ \dots \circ e_1$.

Proof. Each event restricts or quotients the region; neither increases distinguishable futures. \square

18.1 Discrete Action Functional

Define $\mathcal{O}_t = \mu(\Omega_t)$, $L_t = \mathcal{O}_{t-1} - \mathcal{O}_t$, and action $S[H] = \sum_{t=1}^n L_t$.

Theorem 15 (Monotone Action Under Restriction). *If the history contains no collapse events, then $S[H] \geq 0$.*

Proof. Each non-collapse event reduces optionality, giving $L_t \geq 0$. \square

18.2 The Manifold as the Locus of Lawful History

Define

$$M := \{x_H \mid H \text{ is a valid irreversible history}\}.$$

Proposition 11. *Under finite option-spaces and bounded event generators, M admits the structure of a Whitney-stratified manifold embedded in \mathbb{R}^n .*

Proof. Smooth parameter dependence of events generates smooth strata; collapse events yield lower-dimensional quotient strata. \square

Tangent alignment is therefore equivalent to history consistency.

18.3 Minimal Commitment and RSVP Fields as Continuum Limit

The discrete action defines a Morse function $S(x_H) := S[H]$ on M , and gradient flow $\dot{x} = -\nabla_M S(x)$ corresponds in the continuum limit to tangent-constrained descent. The

RSVP fields (Φ, v, S) over (X, g) arise as the smooth relaxation of accumulated irreversible pruning, with embedding evolution

$$\frac{dz_i}{dt} = -\alpha \nabla_X \Phi(z_i) + \beta v(z_i) - \gamma \nabla_X S(z_i).$$

19 Continuum Limit of Irreversible History

We now make precise the relationship between discrete irreversible event histories and continuous RSVP field dynamics.

Let $H_n = e_n \circ \dots \circ e_1$ be a history with discrete action

$$S[H_n] = \sum_{t=1}^n L_t, \quad L_t = \mathcal{O}_{t-1} - \mathcal{O}_t.$$

Define the semantic embedding $x_n := x_{H_n} \in M$ and suppose event generators depend smoothly on parameters $\theta \in \mathbb{R}^k$. Assume bounded event magnitudes and uniform scaling $\Delta t \rightarrow 0$.

Theorem 16 (Continuum Limit of Irreversible History). *Under bounded event generators and finite option-spaces, the rescaled discrete action functional converges to a smooth Morse potential*

$$S(x) = \lim_{\Delta t \rightarrow 0} S[H_n],$$

and discrete update dynamics converge to the gradient flow

$$\dot{x} = -\nabla_M S(x).$$

Proof. Under smooth parameter dependence, the discrete action increments L_t define a Riemann sum approximating an integral functional. Boundedness ensures uniform convergence. The Euler-Lagrange equations of the limiting functional produce intrinsic gradient flow. Convergence of difference quotients to the derivative yields the result. \square

Corollary 2 (Emergence of RSVP Fields). *The scalar entropy field $S(x)$ and associated vector transport field arise as the smooth relaxation of accumulated irreversible pruning in option-space.*

Thus the discrete merge-collapse history and the continuous RSVP dynamics are not independent mechanisms but scale-separated descriptions of the same constraint-driven evolution.

20 Integration with a Typed Field–Hypergraph Verification Architecture

20.1 Typed System Integration and Operational Semantics

The global state is $\Sigma = (C, Z, F, U)$. Small-step transitions $\Sigma \xrightarrow{\alpha} \Sigma'$ cover content ingestion, field evolution, and embedding evolution.

Theorem 17 (Embedding Flow Stability). *If the field potential satisfies convexity conditions and embeddings evolve under the gradient law, then $E(Z) = \sum_i V(z_i)$ is non-increasing.*

Proof. $\frac{dE}{dt} = \sum_i \nabla V(z_i) \cdot \frac{dz_i}{dt} \leq 0$ under gradient-descent structure and convexity. \square

20.2 Denotational Interpretation and Categorical Soundness

Interpret C as a category \mathcal{C} and define $\mathcal{R} : \mathcal{C} \rightarrow \mathcal{F}$ mapping $a \mapsto (\Phi(z_a), v(z_a), S(z_a))$.

Proposition 12 (Functorial Consistency). *If merge-collapse operations preserve typed morphisms, then \mathcal{R} is a functor.*

Theorem 18 (Commutativity up to Natural Transformation). *The diagram formed by the hypergraph endofunctor Poly and \mathcal{R} commutes up to natural transformation.*

Proof. Structural deformation corresponds to induced field perturbation via embedding update and density/curvature recomputation. Naturality follows from preservation of morphism typing and embedding projection. \square

21 A Master Theorem of Coherence: Energy, Gluing, and Soundness

21.1 Standing Definitions and Local Lemmas

Let $E(\Sigma) = \int_X V(\Phi, v, S; \rho_Z, \kappa_C) d\mu_g$ with tangent-constrained embedding dynamics, presheaf \mathcal{G} of deterministic derived views, and content category \mathcal{C} with endofunctor Poly and interpretation functor \mathcal{R} .

Lemma 1 (Lyapunov Monotonicity). *Under dissipative tangent-constrained embedding flow, $\frac{d}{dt} E(\Sigma(t)) \leq 0$.*

Proof. Time derivative of E decomposes into squared-norm terms with negative coefficients; tangent constraint ensures intrinsic interpretation. \square

Lemma 2 (Sheaf Gluing from Deterministic Restriction). *If $\mathcal{G}(U)$ is a deterministic function of restricted authoritative state and restriction is functorial, then \mathcal{G} satisfies the sheaf gluing condition.*

Proof. Overlap agreement implies identical restricted data. Uniqueness follows from determinism. \square

Lemma 3 (Typed Quotient Soundness). *If merge-collapse induces a typing-respecting congruence on \mathcal{C} , then \mathcal{C}/\sim is well-defined and $\overline{\text{Poly}}$ is an endofunctor.*

Proof. Congruence ensures composition descends to classes; typing respect gives well-defined source and target. \square

21.2 Master Theorem

Theorem 19 (Master Coherence Theorem). *Under the four conditions — deterministic replay, typed congruence from merge-collapse, dissipative tangent-constrained dynamics, and deterministically computed derived views — the integrated architecture simultaneously satisfies:*

- (i) *Energy monotonicity: $E(\Sigma(t))$ is non-increasing.*
- (ii) *Contextual coherence: overlap-compatible local views glue uniquely.*
- (iii) *Structural soundness: admissible transformations commute with interpretation up to natural transformation; merge-collapse preserves typing and composition.*

These invariants are mutually stable under reset operations implemented as replay-derived reconstruction.

Proof. (i) is Lemma 1. (ii) is Lemma 2. (iii) is Lemma 3 with functoriality of \mathcal{R} and Poly . Naturality follows from structure-preserving construction of all components. Reset stability follows because reset is a replay-derived reconstruction under the same dissipative flow; all lemmas re-apply. \square

Corollary 3 (No-Normal Drift). *Under the Master Coherence Theorem hypotheses, if manifold membership is enforced by intrinsic gradients or projection after each update, embedding evolution exhibits no normal-component drift.*

22 A Compact Invariant Equation for Semantic Coherence

Let $\Sigma_t = (C_t, Z_t, F_t, H_t)$. Define the global evolution operator:

$$\mathcal{U}(\Sigma_t) = (\mathcal{Q}(C_t), \Pi_T Z_t, \text{Flow}(F_t), H_t \cup \{e_{t+1}\}).$$

The unified invariant is:

$$\boxed{\Pi_T \circ \nabla \mathcal{E} = \nabla \mathcal{E} \quad \text{and} \quad \rho_{UV} \circ \mathcal{U} = \mathcal{U} \circ \rho_{UV}.}$$

Theorem 20 (Unified Structural Invariant). *If the compact invariant equation holds, then all five properties of the Master Coherence Theorem hold: energy monotonicity, manifold confinement, categorical composition preservation, unique gluing, and deterministic replay.*

Proof. Tangent-projected gradient descent implies Lyapunov monotonicity (1) and manifold invariance (2). \mathcal{Q} preserves composition (3). Commutation of \mathcal{U} with restriction gives sheaf gluing (4). Totally ordered append-only log gives determinism (5). \square

Semantic evolution is tangent-projected energy descent that commutes with restriction. All stability, coherence, and soundness results follow from this.

23 Whitney Stratification and Derived Critical Loci

23.1 Stratified Semantic Space

Let $X = \bigsqcup_{\alpha \in A} S_\alpha$ be Whitney-stratified. Semantic evolution is

$$\dot{x} = -\Pi_{T_x S_\alpha} \nabla V(x), \quad x \in S_\alpha.$$

Theorem 21 (Stratum Invariance). *If $x(0) \in S_\alpha$ and the projected gradient is used, then $x(t)$ remains in $\overline{S_\alpha}$ until a controlled boundary transition.*

Proof. $\dot{x} \in T_x S_\alpha$ prevents leaving the closure except at singular boundary points. \square

23.2 Derived Critical Locus and Stratified Morse Transitions

The derived critical locus is the homotopy fiber of $d\mathcal{A}$ over zero.

Proposition 13 (Derived Stability Criterion). *A semantic configuration is structurally stable if and only if its derived critical locus has trivial higher homology.*

Proof. Trivial higher homology implies non-degenerate Hessian structure within strata and absence of hidden deformation directions. \square

A stratified Morse transition from S_α to S_β corresponds to $X_{c+\epsilon} \simeq X_{c-\epsilon} \cup e^\lambda$ at index λ : the emergence or collapse of conceptual dimensions. The compact invariant generalizes to

$$\Pi_{T_x S_\alpha} \circ \nabla \mathcal{E} = \nabla \mathcal{E}, \quad x \in S_\alpha.$$

Category collapse and paradigm shifts are controlled stratum descents, not noise.

24 Stratified Invariant Equivalence

The compact invariant equation was shown to imply the Master Coherence conditions in the smooth manifold case. We now extend this equivalence to the Whitneystratified setting.

Let

$$X = \bigsqcup_\alpha S_\alpha$$

be a Whitneystratified semantic space. Define the stratified compact invariant

$$\Pi_{T_x S_\alpha} \circ \nabla \mathcal{E} = \nabla \mathcal{E}, \quad \rho_{UV} \circ \mathcal{U} = \mathcal{U} \circ \rho_{UV}.$$

Theorem 22 (Stratified Invariant Equivalence). *On a Whitneystratified semantic space, the stratified compact invariant equation is equivalent to the Master Coherence conditions within each stratum and remains stable under controlled boundary transitions.*

Proof. Within each stratum S_α , smooth tangent bundles exist and the smooth equivalence proof applies directly. At boundary points, Whitney conditions ensure compatibility of tangent cones. Since projection is taken relative to $T_x S_\alpha$, flow remains intrinsic to the current stratum until a Morse-type boundary descent occurs. Restriction commutation is unaffected by stratification because contextual maps depend only on embedding data, not ambient smoothness. Therefore the Master Coherence conditions hold piecewise and are preserved across stratified transitions. \square

Corollary 4 (Stratified Coherence). *Semantic bifurcation, category collapse, and conceptual regime shifts are coherent if and only if they correspond to stratified Morse transitions under*

the invariant constraint.

Thus the compact invariant equation governs both smooth evolution and singular transitions. The architecture remains structurally closed under stratification.

25 Operationalizing Manifold-Aware Evaluation (MAE)

The preceding sections establish that coherent evolution requires confinement to the tangent structure of a semantic manifold, smooth or stratified. We now translate this geometric constraint into an evaluative principle for institutional and algorithmic systems. The pathologies described in Section 3 arise when optimization proceeds in ambient directions rather than intrinsic ones. Manifold-Aware Evaluation replaces scalar proxy maximization with alignment-sensitive assessment relative to the manifold M .

25.1 The Manifold Fidelity Score

Let $x \in M$ represent the current institutional or semantic state, and let $\Delta x \in \mathbb{R}^n$ denote a proposed update. Decompose

$$\Delta x = \Pi_{T_x M}(\Delta x) + \Pi_{N_x M}(\Delta x).$$

Definition 9 (Manifold Fidelity Score). *The Manifold Fidelity Score of an update Δx at $x \in M$ is*

$$\mathcal{F}(\Delta x) = \frac{\|\Pi_{T_x M}(\Delta x)\|}{\|\Delta x\|}.$$

This score satisfies $0 \leq \mathcal{F}(\Delta x) \leq 1$. A value of 1 indicates purely intrinsic evolution; values approaching 0 indicate that the proposed change is dominated by normal components and therefore lacks semantic justification relative to the manifold.

In stratified settings, $T_x M$ is replaced by the tangent cone $T_x S_\alpha$ of the stratum containing x , and \mathcal{F} is computed relative to this cone.

25.2 Differential Auditing

Standard metric evaluation examines whether a scalar objective $J(x)$ increases. Manifold-Aware Evaluation instead audits the direction of motion. Let \dot{x} denote the effective institutional update induced by policy or optimization.

Definition 10 (Normal Drift Magnitude).

$$D(x) = \|\Pi_{N_x M}(\dot{x})\|.$$

Persistent nonzero $D(x)$ indicates systematic deviation from intrinsic structure. In this case, the evaluation regime itself is structurally misaligned. Correction requires a Reset Transformation that recomputes fields and embeddings via tangent-constrained relaxation, restoring compatibility with the manifold constraint.

Thus, evaluation becomes directional rather than scalar: the question is not merely whether performance increases, but whether evolution remains intrinsic.

25.3 Curvature-Sensitive Incentives

In regions where the manifold exhibits high curvature, local linear approximations become unreliable. Suppose incentives are defined by a covector field $w(x)$ inducing updates $\Delta x = \epsilon w(x)$. If M is nonlinear in a neighborhood of x , naive linear incentives may generate normal components through second-order effects.

Let $\gamma(t)$ be a geodesic in M . Incentives should be transported along γ via parallel transport to preserve intrinsic alignment. Formally, if ∇ denotes the LeviCivita connection on M , then w should evolve according to

$$\nabla_{\dot{\gamma}} w = 0,$$

ensuring that incentive structure remains intrinsic as the manifold curves.

This curvature-sensitive formulation prevents overshooting the manifold during periods of rapid structural change.

25.4 Optional Information-Geometric Refinement

If M is equipped with a Riemannian metric induced by a Fisher information tensor g_{ij} , the Manifold Fidelity Score may be refined to measure information-theoretic deviation:

$$\mathcal{F}_g(\Delta x) = \frac{\sqrt{g(\Pi_{T_x M} \Delta x, \Pi_{T_x M} \Delta x)}}{\sqrt{g(\Delta x, \Delta x)}}.$$

This formulation weights deviations according to statistical distinguishability, thereby integrating information geometry with manifold-aware auditing.

25.5 A Curvature-Bounded Goodhart Instability Theorem

We now formalize the sense in which linear proxy incentives generically induce normal drift when the semantic constraint set is curved. The result gives a sufficient condition under which proxy-gradient updates exit any prescribed tubular neighborhood of the manifold, even when the proxy appears locally reasonable.

Let $M \subset \mathbb{R}^n$ be a C^2 embedded submanifold of dimension d . Fix $x \in M$. Let $\Pi_x : T_x M \times T_x M \rightarrow N_x M$ denote the second fundamental form of the embedding at x , and let

$$\|\Pi_x\| := \sup_{\substack{u \in T_x M \\ \|u\|=1}} \|\Pi_x(u, u)\|.$$

Let $\tau(M)$ be the reach of M (the largest radius for which nearest-point projection to M is single-valued). For $0 < r < \tau(M)$, denote the tubular neighborhood by

$$\mathcal{T}_r(M) := \{y \in \mathbb{R}^n : \text{dist}(y, M) < r\}.$$

We consider a proxy objective $\tilde{J} : \mathbb{R}^n \rightarrow \mathbb{R}$ and its ambient gradient update

$$x^+ = x + \eta \nabla \tilde{J}(x),$$

with step size $\eta > 0$. Write $g := \nabla \tilde{J}(x)$ and decompose $g = g_T + g_N$ with $g_T \in T_x M$ and $g_N \in N_x M$.

Theorem 23 (Curvature-Induced Goodhart Instability). *Assume $x \in M$ and $0 < r < \tau(M)$. Suppose $g_N = 0$ at x , so the proxy gradient is tangent at the point of evaluation. If $\|g_T\| \neq 0$ and $\|\Pi_x\| > 0$, then there exists $\eta_0 > 0$ such that for all $0 < \eta \leq \eta_0$,*

$$\text{dist}(x + \eta g, M) \geq \frac{\eta^2}{2} \|\Pi_x\| \|g_T\|^2 - O(\eta^3).$$

In particular, for any prescribed $r < \tau(M)$, there exists a step size $\eta_r > 0$ such that for all $\eta \geq \eta_r$ sufficiently small relative to the local geometry, the update leaves the tubular neighborhood:

$$x + \eta g \notin \mathcal{T}_r(M).$$

Thus, even a proxy whose gradient is tangent at evaluation points produces systematic normal deviation at second order whenever the manifold is curved.

Proof. Choose an orthonormal frame adapted to M at x so that, in a neighborhood of x , the manifold can be represented as the graph of a C^2 map $h : T_x M \rightarrow N_x M$ with $h(0) = 0$ and

$Dh(0) = 0$. In these coordinates, points of M near x have the form $u + h(u)$ with $u \in T_x M$. The Taylor expansion gives

$$h(u) = \frac{1}{2} \text{II}_x(u, u) + O(\|u\|^3).$$

Because $g_N = 0$, the ambient update is $x + \eta g = x + \eta g_T$. Its nearest point on M is $x + \eta g_T + h(\eta g_T)$ up to higher-order terms, hence the normal displacement relative to M is approximately $h(\eta g_T)$. Therefore

$$\text{dist}(x + \eta g, M) \geq \|h(\eta g_T)\| = \frac{\eta^2}{2} \|\text{II}_x(g_T, g_T)\| + O(\eta^3) \geq \frac{\eta^2}{2} \|\text{II}_x\| \|g_T\|^2 - O(\eta^3).$$

For any fixed $r < \tau(M)$, choose η_r small enough that the $O(\eta^3)$ term is dominated and the lower bound exceeds r , which implies $x + \eta g \notin \mathcal{T}_r(M)$. The reach condition ensures the tubular neighborhood behaves regularly and that the distance function is well-defined for sufficiently small steps. This proves the claim. \square

Remark 2 (Interpretation). The theorem isolates a purely geometric mechanism for proxy failure. Even when a proxy gradient is locally tangent, curvature couples tangent motion into normal deviation at second order. This is the differential geometric analogue of Goodhart phenomena: optimization along an ambient proxy systematically exits the constraint set unless the update is actively corrected by projection or by an intrinsic transport rule.

Corollary 5 (Projection or Intrinsic Transport is Necessary). *If an evaluation regime applies repeated ambient proxy-gradient steps without projection or intrinsic correction, then for curved manifolds ($\|\text{II}\| > 0$) there exists a time horizon after which the trajectory exits every fixed tubular neighborhood of M . Conversely, enforcing $x^+ := \pi_M(x + \eta \nabla \tilde{J}(x))$ or transporting incentives intrinsically along M removes the curvature-induced normal component to leading order.*

26 Epistemic Interpretation

The preceding constructions admit a systematic epistemic interpretation. The geometric, dynamical, categorical, and operational layers do not represent distinct explanatory frameworks, but alternative formal expressions of a single constraint governing coherent inference.

At the geometric level, constraint is realized as membership in a manifold $M \subset \mathbb{R}^n$ representing the locus of lawful configurations. The tangent–normal decomposition formalizes the distinction between intrinsic variation and extrinsic noise. Epistemic reliability requires that updates be confined to the tangent bundle $T_x M$, thereby preserving manifold invariance.

At the dynamical level, stability is expressed through gradient descent on a Morse-type potential defined intrinsically on M . Convergence to attractors corresponds to the stabilization of interpretation, while saddle structures encode structured transitions between regimes. Stratified extensions generalize this picture to spaces with controlled singularities.

At the categorical level, coherence across perspectives is modeled by the sheaf condition. Local sections defined on overlapping contexts must agree under restriction and glue uniquely to a global section. Update operators that commute with restriction maps preserve this condition and thereby sustain intersubjective compatibility.

At the operational level, safety is expressed through tangent preservation, energy monotonicity, and reset invariance under deterministic replay. These properties ensure that the authoritative state of the system evolves intrinsically, without introducing extraneous degrees of freedom.

The philosophical injunction that explanation must not exceed the structure available to it admits a precise mathematical formulation. It is equivalent to the requirement that

$$\text{Proj}_{N_x M}(\Delta x) = 0$$

for all admissible updates Δx . Geometric confinement, variational descent, categorical gluing, and operational determinism are thus not independent doctrines but mutually reinforcing manifestations of a single structural invariant. Each formal layer re-expresses the same epistemic constraint in the language appropriate to its domain.

27 Conclusion

We began with a philosophical observation about explanation and its limits. In high-dimensional spaces of representation and action, lawful structure occupies a constrained subset. The central claim of this work has been that coherent cognition, stable institutions, and reliable generative systems all depend upon respecting that constraint. Explanation must remain tangent to the manifold of lawful structure; hallucination is the geometric consequence of drifting into normal directions.

From this starting point we derived a precise mathematical framework. Semantic states were modeled as points on a smooth or Whitney-stratified manifold. Meaningful updates were characterized as tangent-constrained gradient flows. Cognitive dynamics were formalized as Morse descent, and conceptual phase transitions were captured as stratified bifurcations governed by changes in homotopy type. Intersubjective stability was expressed sheaf-theoretically. Optimization pathologies were shown to arise when gradients are taken

orthogonal to intrinsic structure.

These geometric and categorical constructions were integrated into an operational architecture. Irreversible merge-collapse computation provided a discrete substrate. Deterministic replay established an authoritative causal order. Typed hypergraph structure supplied compositional content semantics. RSVP field dynamics supplied a variational smoothing mechanism. Reset operations were interpreted as controlled reconfiguration restoring stationarity.

The Master Coherence Theorem establishes that the geometric, variational, categorical, and contextual layers of the framework are not independent assumptions but mutually entailed consequences of a unified structural constraint. Under tangent-constrained evolution on the semantic manifold, dissipative descent of a Lyapunov energy functional, type-preserving merge-collapse operations on the hypergraph substrate, deterministic reconstruction via append-only event replay, and update operators that commute with contextual restriction maps, the integrated system satisfies four simultaneous invariants: geometric confinement to the manifold of lawful structure, monotonic decrease of the associated energy functional, preservation of categorical composition under admissible transformations, and satisfaction of the sheaf gluing condition across overlapping contexts. The compact invariant equation demonstrates that these properties arise from a single commutation principle: projected gradient evolution must remain intrinsic to the manifold and compatible with restriction.

The stratified extension generalizes this result beyond smooth settings. Discontinuities, regime shifts, and categorical boundaries are modeled as Whitney-stratified transitions governed by stratified Morse dynamics. In this setting, coherence is preserved provided that evolution remains tangent within each stratum and transitions occur through controlled boundary descents. Derived critical loci furnish a homotopical characterization of structural stability, identifying latent degeneracies and measuring the fragility of semantic configurations through higher homological structure.

Within this unified perspective, the semantic manifold specifies the locus of lawful configurations admissible under the history of the system. The event log provides the irreversible construction of that locus through monotone restriction of option-space. The RSVP Lagrangian governs the large-scale relaxation and smoothing of embedding and field dynamics. The typed hypergraph samples and discretizes the manifolds local structure. The sheaf-theoretic framework ensures compatibility of local representations across intersecting contexts. Deterministic replay guarantees reproducibility and invariance of authoritative state under identical histories.

Taken together, these components yield a single structural conclusion. Semantic evolution is coherent precisely when it remains intrinsic to the constraint manifold and compatible

across contextual restrictions. Instability corresponds to excitation of normal directions or to violation of gluing conditions. Meaning is therefore not treated as an emergent byproduct of unrestricted computation, but as a consequence of constrained, history-sensitive evolution governed by geometric and categorical invariants.

References

- [1] H. Whitney, Tangents to an analytic variety, *Annals of Mathematics*, 81(3):496–549, 1965.
- [2] J. Milnor, *Morse Theory*, Princeton University Press, 1963.
- [3] V. Guillemin and A. Pollack, *Differential Topology*, Prentice–Hall, 1974.
- [4] J. M. Lee, *Introduction to Smooth Manifolds*, 2nd ed., Springer, 2013.
- [5] R. Bott and L. W. Tu, *Differential Forms in Algebraic Topology*, Springer, 1982.
- [6] M. Kashiwara and P. Schapira, *Sheaves on Manifolds*, Springer, 1990.
- [7] S. Mac Lane, *Categories for the Working Mathematician*, 2nd ed., Springer, 1998.
- [8] S. Awodey, *Category Theory*, 2nd ed., Oxford University Press, 2010.
- [9] D. I. Spivak, *Category Theory for the Sciences*, MIT Press, 2014.
- [10] J.-L. Verdier, Des catégories dérivées des catégories abéliennes, *Astérisque*, 239, 1972.
- [11] J. Lurie, *Higher Topos Theory*, Princeton University Press, 2009.
- [12] I. M. Gelfand and S. V. Fomin, *Calculus of Variations*, Prentice–Hall, 1963.
- [13] V. I. Arnold, *Mathematical Methods of Classical Mechanics*, 2nd ed., Springer, 1989.
- [14] S. Smale, On gradient dynamical systems, *Annals of Mathematics*, 74(1):199–206, 1961.
- [15] C. A. E. Goodhart, Problems of monetary management: The U.K. experience, in *Papers in Monetary Economics*, Reserve Bank of Australia, 1975.
- [16] D. T. Campbell, Assessing the impact of planned social change, *Evaluation and Program Planning*, 2(1):67–90, 1979.
- [17] P. W. Anderson, More is different, *Science*, 177(4047):393–396, 1972.
- [18] M. W. Hirsch, *Differential Topology*, Springer, 1994.
- [19] G. E. Bredon, *Sheaf Theory*, 2nd ed., Springer, 1997.
- [20] J. R. Munkres, *Topology*, 2nd ed., Prentice Hall, 2000.